

Linear Models: CEF and Best Linear Predictor

Robert Gulotty

University of Chicago

March 18, 2026

Populations, Projections, and Structure

- Today we study relationships between random variables Y and $X = (X_1, \dots, X_k)$ in the population.
- The **Conditional Expectation Function (CEF)**: $m(X) = \mathbb{E}[Y \mid X]$.
A potentially non-linear expectation of Y based on X .
- The **Best Linear Predictor (BLP)**: $\beta^* = \arg \min_b \mathbb{E}[(Y - X'b)^2]$.
The linear projection of Y onto X in the population.
- Later we will use algebra to fit a line in the sample and use probability theory to estimate the BLP.
- Structural (exogeneity) assumptions determine whether the BLP coincides with the CEF.

Properties of Conditional Expectation

■ $m(X) \equiv \mathbb{E}[Y | X]$ (Definition)

* Note: $\mathbb{E}[Y|X]$ is a random variable (a function of X), not a number!

■ $\mathbb{E}[\mathbb{E}[Y | X] | X] = \mathbb{E}[Y | X]$ (Idempotence)

■ $\mathbb{E}[\mathbb{E}[Y | X]] = \mathbb{E}[Y]$ (Law of Iterated Expectations)

■ $m(X) = \arg \min_g \mathbb{E}[(Y - g(X))^2]$ (Best predictor of Y using X)

CEF Error and Mean Independence

- Define the CEF error: $e \equiv Y - m(X)$, so

$$Y = Y + (m(X) - m(X)) = m(X) + (Y - m(X)) = m(X) + e$$

- The CEF error has conditional mean zero (by construction):

$$\mathbb{E}[e|X] = \mathbb{E}[Y - m(X)|X] = m(X) - m(X) = 0$$

- And unconditional mean zero (by LIE):

$$\mathbb{E}[e] = \mathbb{E}[\mathbb{E}[e|X]] = 0$$

- $\mathbb{E}[e|X] = 0$ is called *mean independence*.

Hansen Theorem 2.4: Properties of the CEF Error e

- 1 $\mathbb{E}[e|X] = 0$
- 2 $\mathbb{E}[e] = 0$
- 3 If the r th moment of Y exists, the r th moment of e exists. [Technical]
- 4 For *any measurable function* $h(X)$ where the expected product with e exists $[\mathbb{E}[|h(X)e|] < \infty]$, then, $\mathbb{E}[h(X)e] = 0$

Key Point

These are properties, not assumptions; they follow from LIE and the definition of the CEF: *these properties are why we use the CEF as a mathematical tool.*

Why Covariance Matters for Decomposition

Suppose we decompose a random variable into two parts: $Y = S + N$ (“signal” + “noise”).

The variance of the sum is:

$$\text{Var}(Y) = \text{Var}(S) + \text{Var}(N) + 2 \text{Cov}(S, N)$$

- If $\text{Cov}(S, N) \neq 0$, there is no clean separation: signal and noise share variation, and we cannot attribute variance to one or the other.
- If $\text{Cov}(S, N) = 0$, we get a **clean decomposition**:

$$\text{Var}(Y) = \text{Var}(S) + \text{Var}(N)$$

Question: Is there a function of X that decomposes Y into uncorrelated signal and noise?

Covariance of CEF and Its Error

Yes—the CEF does exactly this. Since e is a random variable, compute its covariance with $m(X)$:

$$\mathbb{E}[e m(X)] = \mathbb{E}[\mathbb{E}[e m(X)|X]] = \mathbb{E}[m(X) \mathbb{E}[e|X]] = \mathbb{E}[m(X) \cdot 0] = 0$$

Because $m(X)$ is a measurable function of X , Theorem 2.4.4 applies.

$$\text{Cov}(e, m(X)) = \mathbb{E}[e m(X)] - \mathbb{E}[e] \mathbb{E}[m(X)] = 0 - 0 \cdot \mathbb{E}[m(X)] = 0$$

So the **CEF decomposition** gives us exactly the clean split we wanted:

$$\text{Var}(Y) = \text{Var}(m(X)) + \text{Var}(e)$$


Why Orthogonality? The Projection Interpretation

This is not a mathematical coincidence. It follows from what the CEF *is*:

- The CEF minimizes $\mathbb{E}[(Y - g(X))^2]$ over **all** functions of X (Thm 2.7).
- In the language of Hilbert spaces, $m(X)$ is the **orthogonal projection** of Y onto the space of all square-integrable functions of X .
- A defining property of orthogonal projection: **the residual is perpendicular to the space you project onto**.
- Since $m(X)$, X , X^2 , $\sin(X)$, $\dots$ are all in that space, the error e must be uncorrelated with *every* function of X .

The CEF decomposition $\text{Var}(Y) = \text{Var}(m(X)) + \text{Var}(e)$ is a Pythagorean theorem: it holds because signal and noise are orthogonal.

Visualizing Projection in L^2 Space



The CEF is the closest point in $L^2(X)$ to Y . The residual e is orthogonal to the subspace.

A Note on the Geometry

The diagram on the previous slide is an *analogy*, not a literal picture. Random variables do not live in $\mathbb{R}^3$ —they live in an abstract **Hilbert space** L^2 .

Euclidean $\mathbb{R}^n$		Hilbert space L^2
Vector $\mathbf{v}$	$\longleftrightarrow$	Random variable Y
Dot product $\mathbf{u} \cdot \mathbf{v}$	$\longleftrightarrow$	$\mathbb{E}[UV]$ (inner product)
Length $\ \mathbf{v}\ ^2$	$\longleftrightarrow$	$\mathbb{E}[Y^2]$ (second moment)
Perpendicularity $\mathbf{u} \cdot \mathbf{v} = 0$	$\longleftrightarrow$	Orthogonality $\mathbb{E}[UV] = 0$
Projection onto subspace	$\longleftrightarrow$	Conditional expectation

You do *not* need to know Hilbert space theory for this course. The takeaway is just the intuition: **projection $\Rightarrow$ orthogonal residual $\Rightarrow$ clean variance decomposition.**

How Much Does the CEF Leave Unexplained?

The CEF $m(X)$ captures the systematic part of Y . The error $e = Y - m(X)$ is everything else.

Unconditional error variance:

$$\sigma^2 \equiv \text{Var}[e] = \mathbb{E}[e^2]$$

The average prediction error across all values of X .

Conditional error variance:

$$\sigma^2(x) = \text{Var}[e|X = x] = \mathbb{E}[e^2|X = x]$$

How uncertain is our prediction *for a specific* $X = x$? Averaging: $\sigma^2 = \mathbb{E}[\sigma^2(X)]$.

Example: Predict vote share from incumbency. $\sigma^2(D=1)$ may differ from $\sigma^2(D=0)$ —incumbents might have more predictable outcomes than challengers.

Mean-Variance Representation

Define $u = \frac{e}{\sigma(X)}$, so $\mathbb{E}[u|X] = 0$ and $\text{Var}[u|X] = 1$. Then:

$$Y = m(X) + \sigma(X) u$$

X matters in **two ways**:

- Through the **mean**: $m(X)$ shifts the expected outcome.
- Through the **variance**: $\sigma(X)$ scales the noise.

If $\sigma(X)$ is constant across all X : **homoskedastic** errors.

If $\sigma(X)$ varies with X : **heteroskedastic** errors.

Example: In a model of household income, $\sigma(\text{Education})$ tends to increase with education—PhDs can be professors or hedge fund managers, while high-school graduates have a narrower income range.

Hansen Theorem 2.6

Theorem 2.6 (More Variables Reduce Error Variance)

Let $e_1 = Y - m(X_1)$ and $e_{12} = Y - m(X_1, X_2)$. Then:

$$\text{Var}[Y] \geq \text{Var}[e_1] \geq \text{Var}[e_{12}]$$

You will always explain (weakly) more of the variance with more variables. This is why R^2 is a poor measure of model fit.

Proof strategy: Show $\text{Var}[m(X_1)] \leq \text{Var}[m(X_1, X_2)]$ —the finer CEF has (weakly) larger variance. Then the **CEF decomposition** $\text{Var}[Y] = \text{Var}[m(X)] + \mathbb{E}[\text{Var}[Y|X]]$ converts this into the result about error variances.

Proof of Theorem 2.6: Setup

Goal: Show $\mathbb{E}[m(X_1)^2] \leq \mathbb{E}[m(X_1, X_2)^2]$, where $m(\cdot)$ denotes the relevant CEF.

Step 1. By LIE, the coarser CEF is a conditional expectation of the finer one:

$$m(X_1) = \mathbb{E}[Y|X_1] = \mathbb{E}[\mathbb{E}[Y|X_1, X_2] \mid X_1] = \mathbb{E}[m(X_1, X_2) \mid X_1]$$

Step 2. Jensen's inequality: for any convex function φ (such as $z \mapsto z^2$),

$$\varphi(\mathbb{E}[Z|X_1]) \leq \mathbb{E}[\varphi(Z) \mid X_1]$$

Apply with $Z = m(X_1, X_2)$ and $\varphi(z) = z^2$:

$$\left(\mathbb{E}[m(X_1, X_2) \mid X_1]\right)^2 \leq \mathbb{E}[m(X_1, X_2)^2 \mid X_1]$$

The left side is $m(X_1)^2$ by Step 1, so:

$$m(X_1)^2 \leq \mathbb{E}[m(X_1, X_2)^2 \mid X_1]$$

Proof of Theorem 2.6: Conclusion

Step 3. Take unconditional expectations of both sides:

$$\mathbb{E}[m(X_1)^2] \leq \mathbb{E}[m(X_1, X_2)^2]$$

Step 4. Convert to variances. Recall $\text{Var}[Z] = \mathbb{E}[Z^2] - (\mathbb{E}[Z])^2$, and by LIE:

$$\mathbb{E}[m(X_1)] = \mathbb{E}[m(X_1, X_2)] = \mathbb{E}[Y]$$

so $(\mathbb{E}[Z])^2$ is the same for both. Subtracting it from the inequality above:

$$\text{Var}[m(X_1)] \leq \text{Var}[m(X_1, X_2)]$$

Step 5. Since $\text{Var}[Y] = \text{Var}[m] + \text{Var}[e]$ (by the CEF decomposition), larger $\text{Var}[m]$ means smaller $\text{Var}[e]$:

$$\text{Var}[e_1] \geq \text{Var}[e_{12}] \quad \square$$

Hansen Theorem 2.7: CEF as Best Predictor

Theorem 2.7 (CEF Is the Best Predictor)

If $\mathbb{E}[Y^2] < \infty$, then for any measurable function $g(X)$: $\mathbb{E}[(Y - m(X))^2] \leq \mathbb{E}[(Y - g(X))^2]$. The CEF $m(X) = \mathbb{E}[Y|X]$ minimizes mean squared prediction error.

Proof:

$$\begin{aligned}
 \mathbb{E}[(Y - g(X))^2] &= \mathbb{E}[(e + m(X) - g(X))^2] && (Y = e + m(X)) \\
 &= \mathbb{E}[e^2] + 2\mathbb{E}[e(m(X) - g(X))] + \mathbb{E}[(m(X) - g(X))^2] \\
 &= \mathbb{E}[e^2] + 0 + \mathbb{E}[(m(X) - g(X))^2] && (\text{by thm 2.4.4}) \\
 &\geq \mathbb{E}[e^2] = \mathbb{E}[(Y - m(X))^2] && (\text{squares} \geq 0)
 \end{aligned}$$

Implications of Theorem 2.7

- Regardless of the distribution of X and Y , the best predictor is the CEF $m(X)$.
 - For example: Consider the intercept only model $\mu = \mathbb{E}[Y]$.
 - The best predictor for Y , among constants, is μ .
- Since the CEF is the best predictor, anything else must be identical or worse.
- We don't know what the joint distribution is, so we don't generally know the CEF, we have to use an estimator.

Linear CEF and Best Prediction

If $m(X)$ is *linear* in $X = (X_1, \dots, X_k, 1)'$, we can write $m(X) = X'\beta$. The **linear regression model** is:

$$Y = X'\beta + e, \quad \mathbb{E}[e|X] = 0$$

If homoskedastic: $\mathbb{E}[e^2|X] = \sigma^2$.

Best prediction property:

- If the CEF is linear, the solution to $\beta = \arg \min_b \mathbb{E}[(Y - X'b)^2]$ gives the right formula for β .
- If the CEF is nonlinear, the solution gives the best linear approximation to the true CEF.

From CEF to Linear Predictors

- The CEF provides: summarization, MMSE prediction, and decomposition.
- But the CEF requires knowing the conditional distribution $f_{Y|X}(y|x)$.
 - Exception: CEF is always linear if X is discrete and all interactions are included.
- We will instead approximate the CEF using a **Best Linear Predictor** (BLP).
- Later we will show that OLS estimates are optimal **if** the CEF is linear and the residuals are homoskedastic.

The CEF in Political Science

- **Voter turnout and age:** $\mathbb{E}[\text{Turnout}|\text{Age}]$ is nonlinear—turnout rises steeply from 18–30, plateaus in middle age, and rises again among retirees.
- **Income and party ID:** $\mathbb{E}[\text{Income}|\text{Party}]$ differs across categories—the CEF is just group means when X is discrete.
- **Conflict and GDP:** $\mathbb{E}[\text{Conflict}|\text{GDP}]$ may be highly nonlinear, with sharp thresholds at low GDP.

Takeaway: The CEF is the target. But when the relationship is complex, we approximate it with a linear predictor and interpret accordingly.

What Have We Learned About the CEF?

CEF Summary

- 1 $m(X) = \mathbb{E}[Y|X]$ is the **best predictor** of Y given X (Thm 2.7).
- 2 The CEF error $e = Y - m(X)$ satisfies $\mathbb{E}[e|X] = 0$ **by construction**—not an assumption (Thm 2.4).
- 3 Signal and noise are uncorrelated: $\text{Cov}(m(X), e) = 0$.
- 4 More variables $\Rightarrow$ (weakly) smaller error variance (Thm 2.6).
- 5 $Y = m(X) + \sigma(X)u$: X shapes both the mean *and* the variance.

Next: We approximate $m(X)$ with a *linear* function—the Best Linear Predictor.

(Best) Linear Predictors (bivariate case)

- Suppose we want a predictor $\mathbb{E}^*(Y|X)$ that is **linear**:

$$f(X) = a + bX$$

Our standard for prediction is to minimize the mean-square error (best):

- Whereas the CEF might be infinitely complex, the BLP is characterized just by two numbers, a and b .

Choose a and b to minimize

$$M = \mathbb{E}[(Y - (a + bX))^2]$$

Solving for Best Linear Predictor (a and b)

FOC w.r.t. a :

$$0 = -2\mathbb{E}[Y] + 2\mathbb{E}[a + bX] \implies a = \mathbb{E}[Y] - b\mathbb{E}[X]$$

FOC w.r.t. b :

$$\begin{aligned} 0 &= -2\mathbb{E}[YX] + 2\mathbb{E}[(a + bX)(X)] \\ \mathbb{E}[YX] &= [\mathbb{E}[Y] - b\mathbb{E}[X]]\mathbb{E}[X] + b\mathbb{E}[X^2] \\ \mathbb{E}[YX] - \mathbb{E}[Y]\mathbb{E}[X] &= b[\mathbb{E}[X^2] - \mathbb{E}[X]^2] \\ b &= \frac{\text{Cov}(X, Y)}{\text{Var}(X)} \end{aligned}$$

Best Linear Predictor

The best linear predictor is the population linear projection:

$$\mathbb{E}^*[Y|X] = \beta_0 + \beta_1 X$$

$$Y = \beta_0 + \beta_1 X + e$$

- Parameters (in Greek)

- The slope $\beta = \frac{\text{Cov}(X, Y)}{\text{Var}(X)}$

- The intercept $\alpha = \mathbb{E}(Y) - \beta \mathbb{E}(X)$

- $\mathcal{P}(Y|X) = \mathbb{E}[Y] + \frac{\text{Cov}(X, Y)}{\text{Var}(X)} \cdot (X - \mathbb{E}[X])$

Deriving Variance of BLP at Minimized Values

$$\beta = \frac{\text{Cov}(X, Y)}{\text{Var}(X)}$$

$$\beta \text{Var}(X) = \text{Cov}(X, Y)$$

$$\begin{aligned} \mathbb{E}[(Y - (\alpha + \beta X))^2] - \mathbb{E}[Y - (\alpha + \beta X)]^2 &= \text{Var}(Y - (\alpha + \beta X)) \\ &= \text{Var}(Y - \beta X) \\ &= \text{Var}(Y) + \beta^2 \text{Var}(X) - 2\beta \text{Cov}(X, Y) \\ &= \text{Var}(Y) + \beta^2 \text{Var}(X) - 2\beta^2 \text{Var}(X) \\ &= \text{Var}(Y) - \beta^2 \text{Var}(X) \\ \sigma_\epsilon^2 &= \sigma_Y^2 - \beta^2 \sigma_X^2 \end{aligned}$$

Example: Population Linear Projection is not the CEF

Example: $Y = X + X^2$, $X \sim N(0, 1)$, $\mathbb{E}[X] = \mathbb{E}[X^3] = 0$, $\mathbb{E}[X^2] = 1$

True CEF: $m(X) = X + X^2$

Population Linear Projection: $Y = \alpha + \beta X + e$.

$$\alpha = \mathbb{E}[Y] - \beta \mathbb{E}[X] = \mathbb{E}[X + X^2] - \beta \cdot 0 = 0 + 1 - 0$$

$$\beta = \text{Cov}(X, Y) / \text{Var}(X) = \text{Cov}(X, X + X^2) / 1 = 1/1$$

$$\mathcal{P}(Y|X) = 1 + 1 * X$$

$$e = m(X) - \mathcal{P}(Y|X) = X^2 - 1$$

$$\mathbb{E}[eX] = \mathbb{E}[(X^2 - 1)X] = \mathbb{E}[X^3 - X] = 0$$

The BLP error $e = X^2 - 1$ is a function of X , so $\mathbb{E}[e|X] = X^2 - 1 \neq 0$.

$\mathbb{E}[Xe] = 0$ holds, but $\mathbb{E}[e|X] = 0$ fails—the error has a systematic nonlinear pattern in X . This is what “weaker” means: $\mathbb{E}[e|X] = 0 \implies \mathbb{E}[Xe] = 0$, but not vice versa.

CEF vs Linear Predictors / Projection

- CEF ($m(X)$)
 - Best prediction, the prediction that minimizes squared error.
 - Can be non-linear.
 - Requires knowing joint distribution.
 - Prediction Errors are uncorrelated with any function of X . ($E(e|X) = 0$.)
- Linear Predictor/ Projection ($\mathcal{P}(Y|X)$)
 - Draws a line (plane/hyperplane)
 - Requires knowing variances and covariances.
 - Prediction Errors are uncorrelated with X ($E(Xe) = 0$.)

CEF vs BLP: The Big Picture

Why Settle for Linear?

- The CEF is the *best possible* predictor—but we rarely know it.
- The BLP is the best *linear* predictor—and we only need $\text{Cov}(X, Y)$ and $\text{Var}(X)$.
- When the CEF is actually linear (e.g., jointly normal X, Y), $\text{BLP} = \text{CEF}$.
- When the CEF is nonlinear, BLP gives the best linear approximation.

Key insight: $\mathbb{E}[Xe] = 0$ (BLP errors uncorrelated with X) is *weaker* than $\mathbb{E}[e|X] = 0$ (CEF errors mean-independent of X). The BLP projection “uses up” the linear information in X , but may leave nonlinear patterns in the residual.

Why Write $\mathbb{E}[Xe] = 0$?

- Saying “ X is exogenous” is intuitive but vague. Writing $\mathbb{E}[Xe] = 0$ buys you three things:
 - 1 **It makes assumptions precise.** “Exogenous” could mean many things. $\mathbb{E}[Xe] = 0$ says *exactly* what must hold: the prediction error $e = Y - X'\beta$ is uncorrelated with X .
 - 2 **It generalizes.** Swap X for an instrument Z and you get IV: $\mathbb{E}[Ze] = 0$. Add more conditions and you get GMM. Every estimator in this course is defined by *which* moment conditions it solves.
 - 3 **It defines the parameter.** The population β is the value that makes $\mathbb{E}[X(Y - X'\beta)] = 0$. The OLS formula $(X'X)^{-1}X'Y$ solves the sample analog of this condition.

From Theory to Moment Conditions: A Preview

Every empirical model in this course follows the same logic:

- 1 **Substantive claim:** A mechanism, institution, or design feature creates a restriction on the joint distribution.
- 2 **Moment condition:** Translate that claim into $\mathbb{E}[g(W, \theta)] = 0$.
- 3 **Estimator:** Find $\hat{\theta}$ that solves the sample analog.

Example: In a randomized experiment, random assignment means potential outcomes are independent of treatment. This implies $\mathbb{E}[De] = 0$, and OLS consistently estimates the ATE.

Counterexample: If subjects self-select into treatment based on expected benefits, $\mathbb{E}[De] \neq 0$. OLS is inconsistent; you need a *different* moment condition (e.g., from an instrument). **The question is always:** Why should this moment condition hold?

BLP in R: Bivariate Case

```
1 # Population formula:  $\beta = \text{Cov}(X,Y) / \text{Var}(X)$ 
2 beta_formula <- cov(x, y) / var(x)
3 alpha_formula <- mean(y) - beta_formula * mean(x)
4
5 # OLS estimates the same thing from a sample
6 mod <- lm(y ~ x)
7 coef(mod) # compare to c(alpha_formula, beta_formula)
```

From Bivariate to Multivariate Linear Prediction

- So far we derived the BLP for a single regressor: $\beta = \text{Cov}(X, Y)/\text{Var}(X)$.
- In practice, we want to predict Y using *many* variables: $X = (X_1, X_2, \dots, X_k, 1)'$.
- The algebra generalizes: scalar variances and covariances become **matrices**.

Definition (Mean Square Prediction Error)

The **MSPE** of a linear predictor $X'\beta$ is:

$$S(\beta) = \mathbb{E}[(Y - X'\beta)^2]$$

*This is the expected squared distance between Y and our prediction. The **best linear predictor** minimizes $S(\beta)$ over all $\beta \in \mathbb{R}^k$.*

Regularity Conditions

- Given Y and $X = \begin{pmatrix} X_1 \\ X_2 \\ X_3 \\ 1 \end{pmatrix}$, we need three conditions for the BLP to exist:

1 $\mathbb{E}[Y^2] < \infty$

(finite second moment of Y)

2 $\mathbb{E}\|X\|^2 < \infty$

(finite second moments of regressors)

3 $\mathbf{Q}_{XX} = \mathbb{E}[XX']$ is positive definite

(no perfect collinearity)

- Conditions 1–2 ensure $S(\beta)$ is finite. Condition 3 ensures a *unique* minimizer.
- If $\mathbf{Q}_{XX}$ is not positive definite, some linear combination of regressors is constant—and the coefficients cannot be separately identified.

Deriving the Linear Projection Coefficient β

$$\begin{aligned} S(\beta) &= \mathbb{E}[(Y - X'\beta)(Y - X'\beta)] \\ &= \mathbb{E}[Y^2] - 2\beta'\mathbb{E}[XY] + \beta'\mathbb{E}[XX']\beta \end{aligned}$$

$$0 = \frac{\partial}{\partial \beta} S(\beta) = -2\mathbb{E}[XY] + 2\mathbb{E}[XX']\beta$$

$$\mathbf{Q}_{XY} = \mathbf{Q}_{XX}\beta$$

$$\mathbf{Q}_{XX}^{-1}\mathbf{Q}_{XY} = \beta$$

$$\mathbb{E}[XX']^{-1}\mathbb{E}[XY] = \beta$$

Linear Projection and Projection Error

$$\mathcal{P}[Y|X] = X'\beta = X'\mathbb{E}[XX']^{-1}\mathbb{E}[XY]$$

The projection error is $e = Y - X'\beta$, so $Y = X'\beta + e$.

The error is uncorrelated with X :

$$\mathbb{E}[Xe] = \mathbb{E}[XY] - \mathbb{E}[XX']\mathbb{E}[XX']^{-1}\mathbb{E}[XY] = 0$$

$$\begin{pmatrix} \mathbb{E}[X_1 e] \\ \mathbb{E}[X_2 e] \\ \vdots \\ \mathbb{E}[1 * e] \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

This is k equations, the last of which is $\mathbb{E}[e] = 0$ when there is a constant.

What If There Is No Intercept?

If X does not include a constant, the moment conditions are only $\mathbb{E}[X_j e] = 0$ for each regressor—there is no equation forcing $\mathbb{E}[e] = 0$.

Consequences:

- The projection line is forced through the origin: $\mathcal{P}[Y|X] = X'\beta$ with no intercept.
- The error e can have a nonzero mean: $\mathbb{E}[e] \neq 0$ in general.
- The usual decomposition $\text{Var}(Y) = \text{Var}(X'\beta) + \text{Var}(e)$ still holds (since $\mathbb{E}[Xe] = 0$), but $\mathbb{E}[Y] \neq \mathbb{E}[X'\beta]$.

Takeaway: Always include a constant unless you have a substantive reason not to. Omitting it forces the regression through the origin and allows systematic over- or under-prediction on average.

The Design Matrix

$\beta = \mathbb{E}[XX']^{-1}\mathbb{E}[XY]$ exists and is unique if $\mathbf{Q}_{XX} = \mathbb{E}[XX']$ is invertible.

$\mathbf{Q}_{XX}$ is the population second moment matrix of X .

For any non-zero $\alpha \in \mathbb{R}^k$

$$\alpha' \mathbf{Q}_{XX} \alpha = \mathbb{E}[\alpha' XX' \alpha] = \mathbb{E}[(\alpha' X)^2] \geq 0$$

This must be strictly > 0 to be invertible, if not, there is no unique solution to $\mathbf{Q}_{XY} = \mathbf{Q}_{XX}\beta$, and we say β is not identified.

BLP in R: Multiple Regressors

Now that we have the matrix formula $\beta = \mathbb{E}[XX']^{-1}\mathbb{E}[XY]$, we can compute it directly:

```
1 # Build the design matrix (include a column of 1s)
2 X <- cbind(1, x1, x2)
3
4 # Solve (X'X) beta = X'y directly, without inverting
5 beta_matrix <- solve(t(X) %*% X, t(X) %*% y)
6
7 # lm() does the same thing
8 coef(lm(y ~ x1 + x2)) # same answer
```

`solve(A, b)` solves $Ax = b$ via matrix factorization without ever computing A^{-1} . This is more numerically stable than `solve(A) %*% b`, which explicitly inverts A and amplifies floating-point errors.

Application: Linear Probability Model (Binary Y)

$$P(Y = 1|X) = \beta_0 + \beta_1 X$$

- Suppose we have a binary outcome: $Y = \mathbb{I}(\omega \in B)$
- β_0 is the probability that $Y = 1$ if $X = 0$
- β_1 is the change in probability that $Y = 1$ given a one unit change in X .
- The BLP is still the best linear predictor: $\beta = \mathbb{E}[XX']^{-1}\mathbb{E}[XY]$.
- However, it is an approximation to a probability, not itself a probability.

Linear Predictor Error Variance

- Hansen defines $\sigma^2 = \mathbb{E}[e^2]$, $Q_{YY} = \mathbb{E}[Y^2]$ and $\mathbf{Q}_{YX} = \mathbb{E}[YX']$, a $1 \times k$ vector.
- Recall, X and β are $k \times 1$ vectors.

$$\begin{aligned}
 \sigma^2 &= \mathbb{E}[(Y - X'\beta)^2] \\
 &= \mathbb{E}[Y^2] - 2\mathbb{E}[YX']\beta + \beta'\mathbb{E}[XX']\beta \\
 &= Q_{YY} - 2\mathbf{Q}_{YX}\mathbf{Q}_{XX}^{-1}\mathbf{Q}_{XY} + \mathbf{Q}_{YX}\mathbf{Q}_{XX}^{-1}\mathbf{Q}_{XX}\mathbf{Q}_{XX}^{-1}\mathbf{Q}_{XY} \\
 &= Q_{YY} - \mathbf{Q}_{YX}\mathbf{Q}_{XX}^{-1}\mathbf{Q}_{XY} \\
 &\equiv Q_{YY \cdot X}
 \end{aligned}$$

Intercepts and Slopes

- We will often exclude the 1 from X and write the linear projection as:

$$Y = X'\beta + \alpha + e$$

- Demeaning we have a useful reinterpretation of the projection coefficient:

$$Y - \mu_Y = (X - \mu_X)'\beta + e$$

$$\beta = (\mathbb{E}[(X - \mu_X)(X - \mu_X)'])^{-1} \mathbb{E}[(X - \mu_X)(Y - \mu_Y)]$$

$$\beta = \text{var}[X]^{-1} \text{cov}(X, Y)$$

Partitions and Regression Sub-Vectors

- Divide our regressors into $X = \begin{pmatrix} X_1 \\ X_2 \end{pmatrix}$, so the linear projection is:

$$Y = X_1'\beta_1 + X_2'\beta_2 + e, \quad \mathbb{E}[Xe] = 0$$

- We can partition $\mathbf{Q}_{XX}$ and $\mathbf{Q}_{XY}$:

$$\mathbf{Q}_{XX} = \begin{bmatrix} \mathbf{Q}_{11} & \mathbf{Q}_{12} \\ \mathbf{Q}_{21} & \mathbf{Q}_{22} \end{bmatrix} \quad \mathbf{Q}_{XY} = \begin{bmatrix} \mathbf{Q}_{1Y} \\ \mathbf{Q}_{2Y} \end{bmatrix}$$

A Useful Formula

- Recall the error variance was: $Q_{YY.X} \equiv Q_{YY} - Q_{YX} Q_{XX}^{-1} Q_{XY}$
- Analogously $Q_{11.2} \equiv Q_{11} - Q_{12} Q_{22}^{-1} Q_{21}$, and $Q_{22.1} \equiv Q_{22} - Q_{21} Q_{11}^{-1} Q_{12}$, etc.
- In each case, $Q_{a.b}$ is the variation in a not predicted by the linear projection of b .

Partitioned Matrix Inversion and Regression

- By similar reasoning:

$$\mathbf{Q}_{XX}^{-1} = \begin{bmatrix} \mathbf{Q}_{11 \cdot 2}^{-1} & -\mathbf{Q}_{11 \cdot 2}^{-1} \mathbf{Q}_{12} \mathbf{Q}_{22}^{-1} \\ -\mathbf{Q}_{22 \cdot 1}^{-1} \mathbf{Q}_{21} \mathbf{Q}_{11}^{-1} & \mathbf{Q}_{22 \cdot 1}^{-1} \end{bmatrix}$$

- As a result:

$$\begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix} = \begin{bmatrix} \mathbf{Q}_{11 \cdot 2}^{-1} \mathbf{Q}_{1Y \cdot 2} \\ \mathbf{Q}_{22 \cdot 1}^{-1} \mathbf{Q}_{2Y \cdot 1} \end{bmatrix}$$

- We can use $\beta_1 = \mathbf{Q}_{11 \cdot 2}^{-1} \mathbf{Q}_{1Y \cdot 2}$ to understand multivariate regression.

Interpreting Multivariate Regression

- The multivariate projection equation is $Y = X_1'\beta_1 + X_2'\beta_2 + e$.
- Decompose X_1 into the part predicted by X_2 and the part that is not:

$$X_1 = \underbrace{X_2'\gamma_2}_{\text{explained by } X_2} + \underbrace{u_1}_{\text{residual}}$$

- β_1 is the projection of Y onto u_1 —the variation in X_1 that is *unique* to X_1 , not shared with X_2 .
- **Example:** Regress growth on democracy (X_1) and institutions (X_2). β_1 captures the effect of democracy *after removing* what institutions already explain about democracy.

Demonstration: Iterated Projection

Divide X into a single variable in X_1 , with all the other variables in X_2 , so that

$$Y = X_1'\beta_1 + X_2'\beta_2 + e$$

Regress X_1 on X_2 :

$$X_1 = X_2'\gamma_2 + u_1$$

$$\mathbb{E}[X_2 u_1] = 0$$

$$\gamma_2 = \mathbf{Q}_{22}^{-1} \mathbf{Q}_{21}$$

$$\mathbb{E}[u_1^2] = \mathbf{Q}_{11 \cdot 2}$$

$$\mathbb{E}[u_1 Y] = \mathbf{Q}_{1Y} - \mathbf{Q}_{12} \mathbf{Q}_{22}^{-1} \mathbf{Q}_{2Y} = \mathbf{Q}_{1Y \cdot 2}$$

$$\beta_1 = \mathbf{Q}_{11 \cdot 2}^{-1} \mathbf{Q}_{1Y \cdot 2} = \frac{\mathbb{E}[u_1 Y]}{\mathbb{E}[u_1^2]}$$

Omitted Variable Bias

- **Long regression:** $Y = X_1'\beta_1 + X_2'\beta_2 + e$. **Short regression:** $Y = X_1'\gamma_1 + u$.
- Substitute the long into the short:

$$\begin{aligned}\gamma_1 &= (\mathbb{E}[X_1 X_1'])^{-1} \mathbb{E}[X_1 Y] \\ &= (\mathbb{E}[X_1 X_1'])^{-1} \mathbb{E}[X_1 (X_1' \beta_1 + X_2' \beta_2 + e)] = \beta_1 + \underbrace{(\mathbb{E}[X_1 X_1'])^{-1} \mathbb{E}[X_1 X_2']}_{\Gamma_{12}} \beta_2\end{aligned}$$

- Γ_{12} : coefficient from regressing X_2 on X_1 (the **association** between included and omitted).
- Bias = $\Gamma_{12}\beta_2$: product of *how correlated* X_1 and X_2 are and *how much* X_2 matters for Y .

OVB is Not Solved by Adding Variables

- Suppose $Y = X'\beta + e = X'_1\beta_1 + X'_2\beta_2 + X'_3\beta_3 + e$
- Suppose we cannot measure X_3 . Instead, we can choose either:

$$Y = X'_1\gamma_1 + u_1$$

$$Y = X'_1\delta_1 + X'_2\delta_2 + u_2$$

- $\gamma_1 = \beta_1 + \Gamma_{12}\beta_2 + \Gamma_{13}\beta_3$.
- $\delta_1 = \beta_1 + \Gamma_{13.2}\beta_3$.
- Which is better depends on the signs and sizes of these terms!

OVV in Practice: Does Democracy Cause Growth?

- Barro (1996) regresses GDP growth on a democracy index:

$$\text{Growth}_i = \beta_1 \text{Democracy}_i + e_i$$

- But countries that democratize also tend to have stronger property rights, higher education, and colonial histories that affect growth.
- The OVB formula tells us:

$$\gamma_1 = \beta_1 + \underbrace{\Gamma_{12}}_{\text{correlation of democracy with institutions}} \underbrace{\beta_2}_{\text{effect of institutions on growth}}$$

- If institutions are positively correlated with democracy ($\Gamma_{12} > 0$) and positively affect growth ($\beta_2 > 0$), then γ_1 **overstates** the effect of democracy.
- Adding controls can help—but only if they reduce the bias term, not increase it.

Where We Stand

Recap: From Population to Practice

- 1 **CEF** ($m(X) = \mathbb{E}[Y|X]$): Best predictor, but requires the full joint distribution.
- 2 **BLP** ($\mathcal{P}[Y|X] = X'\beta$): Best *linear* predictor, requiring only means, variances, and covariances.
- 3 **Partitioned regression**: Each coefficient captures the effect of a variable *after removing* what other variables predict.
- 4 **OVB**: Omitting a correlated variable biases coefficients; adding variables doesn't always help.

Next: All of this is in the *population*. Starting next lecture, we move to *samples*—using OLS to estimate β and studying when OLS does a good job.